

Integers

- | | | | | |
|---------------|-------|-------|-------|-------|
| | 0x100 | 0x101 | 0x102 | 0x103 |
| Big endian | 01 | 23 | 45 | 67 |
| Little endian | 67 | 45 | 23 | 01 |
- $\lceil x/2^k \rceil$ is given by $(x + (1 \ll k) - 1) \gg k$
 - $x/2^k$ is given by $(x < 0 ? x + (1 \ll k) - 1 : x) \gg k$

Floating Point

- Lacks common `+` rules `a*b*c != a*(b*c), a+b+c != a+(b+c), a*(b+c) != a*b + a*c`
- Contains 3 parts `sign,exp,frac`; $V = (-1)^s \times M \times 2^E$
- 3 parts of float: 1+8+23=32; double: 1+11+52=64
- Exponent E , k -bit `exp` field
- Significand (mantissa) $M - 0 \leq M < 1$ or $1 \leq M < 2$, represented by n -bit `frac` field ($f_{n-1} \dots f_1 f_0$)
- Normalized values
 - `exp` is neither all zeros nor all ones
 - $E = e - \text{Bias}$ where e is the unsigned number in `exp` and $\text{Bias} = 2^{k-1} - 1$
 - `frac` represents $0 \leq f < 1$ with $0.f_{n-1} \dots f_1 f_0$ and $M = 1 + f$

1	10000000	10101010101010101010101
$(-1) \times 2^1 \times 1.101010101010101010101$		
- Denormalized values
 - `exp` all zeros
 - Exponent value is $E = 1 - \text{Bias}$, $M = f$ (no leading 1)

0	00000000	10101010101010101010101
$1 \times 2^{1-01111111} \times 0.101010101010101010101$		
- Special values
 - Exponent field is all ones
 - Fraction field all zeros can represent $\pm\infty$, depending on sign bit
 - Nonzero fraction field is NaN
 - $+\infty = 0, 11111111, 000000000000000000000000$
 - every conditional expression with NaN is 0
 - NaN == NaN is 0
 - every arithmetic expression with NaN is NaN
 - NaN * 0 is NaN
 - forcibly trans `inf` or NaN to `int` gets `int_min`
- Rounding
 - 向更近的地方舍入，如果一样，就向偶舍入。比如下面保留四位：
 - $1.01110010011 \rightarrow 1.0111$ ，末尾小于 0.5
 - $1.01111110011 \rightarrow 1.1000$ ，末尾大于 0.5
 - $1.0101[100\dots] \rightarrow 1.0110$ ，末尾恰是 0.5，奇数向偶数
 - $1.0110[100\dots] \rightarrow 1.0110$ ，末尾恰是 0.5，偶数不变
- Multiplication
 - $(-1)^{s_1} M_1 2^{E_1} \times (-1)^{s_2} M_2 2^{E_2}$
 $= (-1)^{s_1 \oplus s_2} (M_1 \times M_2) 2^{E_1 + E_2}$
 - If $M_1 \times M_2 \geq 2$, $M = M_1 \times M_2 / 2, E = E_1 + E_2 + 1$
 - If E out of range, overflow to `inf`
 - Round M to fit `frac` precision
 - 溢出优先于舍入

Assembly Basics

- “word”=16-bit; “double word”=32-bit (int); “quad words”=64-bit
- 64 位机器下，指针是 8-byte 的
- `mov S, D` has the effect of $S \rightarrow D$

<code>movq Src, Dest</code>	C Analog
<code>movq \$0x4, %rax</code>	<code>tmp = 0x4</code>
<code>movq \$-147, (%rax)</code>	<code>*p = -147</code>
<code>movq %rax, %rdx</code>	<code>tmp2 = tmp1</code>
<code>movq %rax, (%rdx)</code>	<code>*p = tmp</code>
<code>movq (%rax), %rdx</code>	<code>tmp = *p</code>
- `movzbq` 把 8 位 (b) 搬运到 64 位 (q) 寄存器，空出的高位补 0(z)
- `movsbq` 把 8 位 (b) 搬运到 64 位 (q) 寄存器，空出的高位根据符号位 (s)
- `leaq S, D` has the effect of $\&S \rightarrow D$
- `subq S, D` has the effect of $D - S \rightarrow D$
- `salq S, D` has the effect of $D \cdot 2^S \rightarrow D$

Opcodes of arithmetic operations					
<code>addq</code>	<code>+</code>	<code>xorq</code>	<code>⊕</code>	<code>salq</code>	<code>also shlq <<</code>
<code>subq</code>	<code>-</code>	<code>andq</code>	<code>&</code>	<code>sarq</code>	<code>Arithmetic >></code>
<code>imulq</code>	<code>×</code>	<code>orq</code>	<code> </code>	<code>shrq</code>	<code>Logical >></code>
<code>incq</code>	<code>++</code>	<code>decq</code>	<code>--</code>	<code>negq</code>	<code>-</code>
				<code>notq</code>	<code>~</code>

- Logical shift** – Fills left end with zeros (unsigned)
 - Arithmetic shift** – Sign-extends left end (signed)
- Operand Types:**
- Immediate:** $\$Imm \rightarrow \text{value is } Imm$
 - Register:** $r_a \rightarrow \text{value is } R[r_a]$
 - Memory:** $Imm(r_b, r_i, s) \rightarrow M[Imm + R[r_b] + R[r_i] \cdot s]$
 - $Imm(r_b, r_i, s) \rightarrow Imm + R[r_b] + R[r_i] * s$

Operands (3 types)		
Type	Form	Operand value
Immediate	$\$Imm$	Imm
Register	r_a	$R[r_a]$
Memory	$Imm(r_b, r_i, s)$	$M[Imm + R[r_b] + R[r_i] \cdot s]$

16 general purpose registers storing 64-bit values

Type	64-bits	32-bits	16-bits	8-bits
Registers below are caller saved				
Return val	<code>%rax</code>	<code>%eax</code>	<code>%ax</code>	<code>%al</code>
1st arg	<code>%rdi</code>	<code>%edi</code>	<code>%di</code>	<code>%dil</code>
2nd arg	<code>%rsi</code>	<code>%esi</code>	<code>%si</code>	<code>%sil</code>
3rd arg	<code>%rdx</code>	<code>%edx</code>	<code>%dx</code>	<code>%dl</code>
4th arg	<code>%rcx</code>	<code>%ecx</code>	<code>%cx</code>	<code>%cl</code>
5th arg	<code>%r8</code>	<code>%r8d</code>	<code>%r8w</code>	<code>%r8b</code>
6th arg	<code>%r9</code>	<code>%r9d</code>	<code>%r9w</code>	<code>%r9b</code>
Caller	<code>%r10</code>	<code>%r10d</code>	<code>%r10w</code>	<code>%r10b</code>
Caller	<code>%r11</code>	<code>%r11d</code>	<code>%r11w</code>	<code>%r11b</code>

`%rsp` is the top of the stack and `%rbp` is the bottom

Conditional Control

set D and jmp suffixes			
Instruction	Alias.	Cond.	Desc.
<code>-e</code>	<code>-z</code>	ZF	<code>= / 0</code>
<code>-ne</code>	<code>-nz</code>	<code>~ZF</code>	<code>! = / not zero</code>
<code>-s</code>		SF	Neg
<code>-ns</code>		<code>~SF</code>	Nonneg
<code>-g</code>	<code>-nle</code>	<code>~(SF^OF)&~ZF</code>	signed <code>></code>
<code>-ge</code>	<code>-nl</code>	<code>~(SF^OF)</code>	signed <code>>=</code>
<code>-l</code>	<code>-nge</code>	SF^OF	signed <code><</code>
<code>-le</code>	<code>-ng</code>	<code>(SF^OF) ZF</code>	signed <code><=</code>
<code>-a</code>	<code>-nbe</code>	<code>~CF&~ZF</code>	unsigned <code>></code>
<code>-ae</code>	<code>-nb</code>	<code>~CF</code>	unsigned <code>>=</code>
<code>-b</code>	<code>-nae</code>	CF	unsigned <code><</code>
<code>-be</code>	<code>-na</code>	CF ZF	unsigned <code><=</code>

`cmpq B, A, jg L1` equals to if (A > B) goto L1;
`setne %al` sets `%al` to 1 if not equal, 0 otherwise

- Carry flag (CF) 最高位进位，检查无符号操作溢出。
- Zero flag (ZF) 最近的操作：结果为 0。
- Sign flag (SF) 最近的操作：结果为负数。
- Overflow flag (OF) 补码溢出，检查有符号操作溢出。
- `test` 模拟按位 `and`，不保留结果，只改变寄存器。
- `cmp B, A` 模拟减法 `A-B`，不保留结果，只改变寄存器。
- Switch 语句:** 跳转表 `jtab`[索引 `i`] 对应 `x=i` 的地址。
- 编译器通常生成 `leaq + jmp * (%reg)`

Machine Data

- K 字节数据需存放在地址为 K 的倍数位置 (Alignment)
- 基本类型大小：

Size (bytes)	Types
1	<code>char</code>
2	<code>short</code>
4	<code>int, float</code>
8	<code>long, double, char *</code>

- Structs:**
- 每个字段的起始地址需满足自身类型对齐要求
- 字段间可能插入 `padding`
- 结构体总对齐要求 = 所有字段最大对齐
- 结构体总大小 = 最大对齐的整数倍
- Example: `struct {char a; int b; char c;} →`

12 bytes	0	1	2	3	4	5	6	7	8	9	10	11
	a	pad			b			c		pad		

- Nested structs: 内层 struct 按自身对齐计算后整体视为一个字段
- Access field at offset d: `movq d(%rdi), %rax`
- Unions:**
- 所有字段共享内存起始地址（偏移为 0）
- 联合体总大小 = 最大成员大小
- Example: `union {int i; float f; char c[4];}`

occupies 4 bytes total	0	1	2	3
	i (4 bytes)			
	f (4 bytes)			
	c[0]	c[1]	c[2]	c[3]

- Access any member using base address: `movl (%rdi), %eax`
- Used for type punning 双关 or memory optimization
- `union{float f; int i;};` 是用来查看浮点数内部位模式。

SIMD

- `setzero, load / add, hadd / cast, extract / store, cvtss`
- `ps(float), pd(double), epi32(int)`

```
float arraySumVector(const float* values, int N) {
    const int W = 8; int i = 0;
    __m256 sum256 = _mm256_setzero_ps();
    for (; i <= N - W; i += W) {
        __m256 vec = _mm256_loadu_ps(values + i);
        sum256 = _mm256_add_ps(sum256, vec);
    }
    __m128 low = _mm256_castps256_ps128(sum256);
    __m128 high = _mm256_extractf128_ps(sum256, 1);
    __m128 sum = _mm_add_ps(low, high);
    sum = _mm_hadd_ps(sum, sum);
    sum = _mm_hadd_ps(sum, sum);
    float ans = _mm_cvtss_f32(sum);
    for(; i < N; i++) ans += values[i];
    return ans;
}
```

GNU Debugger

- run kill step next stepi continue
- finish: Continue until current function returns
- break <function name> / <file>:<line number> / *(<address>)
- delete <num> clear enable / disable <num>
- info / breakpoints / registers / functions / locals / args / displays
- layout, layout asm, regs, src, update
- print [/FMT] <expr>
- x [/N] [FMT] [SIZE] <addr> x/4xw \$rsp
- see help x in GDB. hex(16),dec(10),oct(8),bin(2).

Linking

- **ELF**. 链接关注: Section, 执行关注: Segment.
- **Symbols** 不包括 Local non-static, 与 linker 无关。
- **Global**: 函数和普通全局变量。
- Global 分为 **Strong**: 函数/初始化的; **Weak**: 未初始化。
- **External**: 定义在其他文件中。
- **Local**: 用 **static** 定义的函数和全局变量。
- **COMMON**: 未初始化的全局变量。
- .bss: 未初始化静态变量, 为 0 的全局和静态变量。
- **Relocation**. PC32 是相对寻址, 32 是绝对寻址。
- offset 是代码中的位置, addend 是跳转偏移量。

Processes

- getpid(), getppid()-parent, exit(int) 永不返回
- fork(void) pid-父, 0-子; 返回两次, 先后顺序不定
- Processes are running, stopped, or terminated
- wait(int *statusp) \cong waitpid(-1, &status, 0)
- sleep, tpauses(), execve()

Signals

- Signals 不是按顺序处理的。
- setpgid(0, 0) 会以当前 PID 作为 ID 创建新进程组
- sigprocmask(int how, sigset_t *set, sigset_t *oldset)
- sigemptyset(sigset_t *set)
- sigfillset(sigset_t *set)
- sigaddset(sigset_t *set, int sig) - add sig
- sigdelset(sigset_t *set, int sig) - delete sig
- sigsuspend(sigset_t *mask)

System I/O

- open, open("foo.txt", O_RDWR | O_CREAT);
- read, read(fd, buf, 100);
- write, write(fd, msg, strlen(msg));
- 每个进程都有唯一的描述符表, 描述符是小整数。
- 操作系统维护所有进程共享的打开文件表
- 每个条目都有文件位置、引用计数和 v-node 表的指针
- 操作系统维护 v-node 表, 其中包含有关每个文件的信息
- fork() 后子进程会继承父进程的文件描述符表
- 父进程和子进程必须关闭文件, 以便内核删除文件表条目
- dup(int oldfd), dup2(int oldfd, int newfd)
- 0-stdin, 1-stdout, 2-stderr, 3... for dup

Tiny Shell

```
void eval_external(struct cmdline_tokens tok, int bg, char *cmdline) {
    pid_t pid; sigset_t mask, mask_old;
    sigemptyset(&mask); sigaddset(&mask, SIGCHLD);
    sigaddset(&mask, SIGINT); sigaddset(&mask, SIGSTP);
    sigprocmask(SIG_BLOCK, &mask, &mask_old);
    if((pid = fork()) == 0) {
        sigprocmask(SIG_SETMASK, &mask_old, NULL); setpgid(0, 0);
        execve(tok.argv[0], tok.argv, environ);
        sio_put("Command not found: %s\n", tok.argv[0]); exit(0);
    } else { addjob(job_list, pid, bg ? BG : FG, cmdline);
        sigprocmask(SIG_SETMASK, &mask_old, NULL);
        if(bg) {
            printf("[%d] (%d) %s\n", pid2jid(pid), pid, cmdline);
        } else { sigset_t mask_zero; sigemptyset(&mask_zero);
            while(fgpid(job_list) == pid) sigsuspend(&mask_zero); } } }

void eval_builtin_bg(struct cmdline_tokens tok) {
    int jid = atoi(tok.argv[1] + 1);
    struct job_t *job = getjobjid(job_list, jid);
    if(job == NULL) { printf("%s: No such job\n", tok.argv[1]); return; }
    pid_t pid = job->pid; kill(pid, SIGCONT); job->state = BG;
    printf("[%d] (%d) %s\n", jid, pid, job->cmdline);
}

void eval_builtin_fg(struct cmdline_tokens tok) {
    int jid = atoi(tok.argv[1] + 1);
    struct job_t *job = getjobjid(job_list, jid);
    if(job == NULL) {
        printf("%s: No such job\n", tok.argv[1]); return;
    }
    pid_t pid = job->pid; kill(pid, SIGCONT);
    job->state = FG; sigset_t mask_zero; sigemptyset(&mask_zero);
    while(pid == fgpid(job_list)) sigsuspend(&mask_zero);
}

int get_kill_pid(char *arg, pid_t *kill_pid, int *is_job_id) {
    if(arg[0] == '%') {
        *is_job_id = 1; long jid = strtol(arg + 1, NULL, 10);
        jid = jid < 0 ? -jid : jid;
        struct job_t *job = getjobjid(job_list, jid);
        if(!job) return -1; *kill_pid = -job->pid;
    } else {
        *is_job_id = 0; long pid = strtol(arg, NULL, 10);
        if(!pid) return -1; *kill_pid = (pid_t)pid; } return 0;
}

void eval_builtin_kill(struct cmdline_tokens tok) {
    if(tok.argv[1] == NULL) {
        printf("kill command requires PID or %%jobid argument\n");
```

```
        return;
    }
    pid_t kill_pid; int is_job_id;
    if(get_kill_pid(tok.argv[1], &kill_pid, &is_job_id) < 0) {
        if(is_job_id) printf("%s: No such job\n", tok.argv[1]);
        else printf("%s: No such process group\n", tok.argv[1]);
        return;
    }
    kill(kill_pid, SIGTERM);
}

void eval_builtin_nohup(char *cmdline) {
    sigset_t mask, mask_old;
    sigemptyset(&mask); sigaddset(&mask, SIGHUP);
    sigprocmask(SIG_BLOCK, &mask, &mask_old);
    eval(cmdline + 6); sigprocmask(SIG_SETMASK, &mask_old, NULL);
}

void eval(char *cmdline) { /*...
    int new_input_fd, new_output_fd, old_input_fd, old_output_fd;
    if(tok.infile != NULL) { old_input_fd = dup(STDIN_FILENO);
        new_input_fd = open(tok.infile, O_RDONLY);
        dup2(new_input_fd, STDIN_FILENO); close(new_input_fd); }
    if(tok.outfile != NULL) { old_output_fd = dup(STDOUT_FILENO);
        new_output_fd = open(tok.outfile, O_WRONLY | O_TRUNC | O_CREAT);
        dup2(new_output_fd, STDOUT_FILENO); close(new_output_fd); }
    switch(tok.builtins) {
        case BUILTIN_NONE: eval_external(tok, bg, cmdline); break;
        /* ... */
    }
    if(tok.infile != NULL) {
        dup2(old_input_fd, STDIN_FILENO); close(old_input_fd);
    }
    if(tok.outfile != NULL) {
        dup2(old_output_fd, STDOUT_FILENO); close(old_output_fd);
    }
    return;
}

void sigchld_handler(int sig) {
    int errno_old = errno;
    pid_t pid; int status; struct job_t *job;
    sigset_t mask, mask_old;
    sigfillset(&mask);
    sigprocmask(SIG_BLOCK, &mask, &mask_old);
    while((pid = waitpid(-1, &status,
        WNOHANG | WUNTRACED | WCONTINUED)) > 0) {
        job = getjobpid(job_list, pid);
        if(WIFEXITED(status)) deletejob(job_list, pid);
        else if(WIFSTOPPED(status)) {
            sio_put("Job [%d] (%d) stopped by signal %d\n",
                pid2jid(pid), pid, WSTOPSIG(status));
            job->state = ST;
        } else if(WIFSIGNALED(status)) {
            sio_put("Job [%d] (%d) terminated by signal %d\n",
                pid2jid(pid), pid, WTERMSIG(status));
            deletejob(job_list, pid);
        } else if(WIFCONTINUED(status)) job->state = BG;
    }
    sigprocmask(SIG_SETMASK, &mask_old, NULL);
    errno = errno_old; return;
}

void sigint_handler(int sig) {
    int errno_old = errno; pid_t pid = fgpid(job_list);
    if(pid > 0) kill(-pid, SIGINT);
    errno = errno_old; return;
}

Optimizaiton

for (int bt = 0; bt < BT; bt += 8) {
    const float *in0 = inp + (bt + 0) * C; /*...*/
    for (int o = 0; o < OC; ++o) {
        float val0 = bias ? bias[o] : 0.0f;
        float val1 = val0, val2 = val0; /*...*/
        float *row_addr = &weight[o * C];
        for (int i = 0; i < C; ++i) {
            float w = row_addr[i];
            float t0 = in0[i]; float t1 = in1[i]; /*...*/
            val0 += t0 * w; val1 += t1 * w; /*...*/
        }
        out[(bt + 0) * OC + o] = val0; /*...*/
    }
}

void advanced_memset(void *s, int c, size_t n) {
    unsigned char *schar = s;
    uintptr_t addr_val = (uintptr_t)schar;
    unsigned char c_byte = (unsigned char)c;
    unsigned long c_long = 0;
    for(size_t i = 0; i < WORD_SIZE; ++i)
        { c_long = (c_long << 8) | c_byte; }
    while(n > 0 && (addr_val % WORD_SIZE != 0)) {
        byte_write(schar, c_byte);
        --n; ++schar; addr_val = (uintptr_t)schar;
    }
    unsigned long *word_ptr = (unsigned long *)schar;
    size_t num_words = n / WORD_SIZE;
    while(num_words >= 4) {
        word_write(word_ptr + 0, c_long); /*...*/
        word_ptr += 4; num_words -= 4;
    }
    while(num_words > 0) {
        word_write(word_ptr, c_long);
        ++word_ptr; --num_words;
    }
    schar = (unsigned char *)word_ptr;
    size_t remaining_bytes = n % WORD_SIZE;
    while(remaining_bytes > 0) {
        byte_write(schar, c_byte);
        --remaining_bytes; ++schar;
    }
}
```